

Executive Summary

Executive Summary: Actor Grounding Layer v0.1

AGL - Precondition Layer for Runtime Reliance

Status	Draft v0.1
Layer	Source-state qualification / pre-review execution grounding
Canonical home	ester-reality-bound
Author	Kotov Ivan
Location	Brussels
Year	2026

Kotov Ivan
Bruxelles, 2026

Contents

1. The Problem	2
2. What AGL Is	2
3. What AGL Qualifies	3
4. Core Principle	3
5. What AGL Is Not	3
6. Relationship to the Existing Stack	4
7. Why the Canonical Home Is ester-reality-bound	4
8. Why This Matters	5
9. Explicit Bridge	5
10. Hidden Bridges	5
11. Earth Paragraph	5
12. Summary	6

1. The Problem

Many AI architectures already distinguish between:

- evidence and narrative,
- lawful and unlawful action,
- review and execution,
- visible state and authoritative state.

That is necessary, but it is no longer sufficient.

A deeper failure appears one step earlier:

the system may validate a dispute, a claim, or an execution path
without first qualifying whether the originating actor or signal is grounded enough to support runtime reliance at all.

In distributed, persistent, or multi-entity environments, this assumption fails quickly. A source may be:

- delegated,
- simulated,
- partially abstracted,
- sensor-corrupted,
- physically detached from the live boundary,
- cognitively degraded,
- physiologically degraded,
- or procedurally valid in form while already unstable in state.

A system may therefore be structurally clean and still act on a source that should never have been allowed to move the runtime forward.

2. What AGL Is

The **Actor Grounding Layer (AGL)** is a bounded architectural layer that asks a narrower and harder question than review:

Is the originating actor, source, or perception path grounded enough in real present execution state to support action, escalation, or review?

AGL does not decide the final dispute. It decides whether a signal may be relied upon strongly enough to:

- proceed,
- open a pre-admissibility hold,
- enter review,
- request escalation,
- or be denied as runtime-relevant input.

AGL is therefore not a commentary layer. It is a **hard gate by default**.

3. What AGL Qualifies

AGL qualifies four source classes:

3.1 Human / operator state

Whether the human source is sufficiently grounded for reliance at the moment of initiation. This may include overload, fatigue, stress, sleep deprivation, or other state degradation.

3.2 Local entity initiation state

Whether the local continuity-bearing entity is in a valid state to initiate, escalate, or authorize action.

3.3 Sensor / perception state

Whether the signal is arriving through a sufficiently grounded perceptual path rather than through corrupted, stale, or operationally detached sensing.

3.4 Delegated / proxy / external origin state

Whether the source remains anchored strongly enough to real execution rather than drifting into abstraction through proxies, delegation chains, or simulated surfaces.

4. Core Principle

The core AGL principle is simple:

Grounding is not only something to assess.

It is something the system may require as a precondition to act at all.

If grounding is missing or materially degraded, the system must not treat normal continuation as the default.

It should instead:

- open a hold,
- narrow privilege,
- require stronger evidence,
- quarantine the path,
- refuse escalation,
- or deny runtime reliance entirely.

This is not pessimism. It is prevention.

5. What AGL Is Not

AGL is **not**:

- a metaphysical detector of authenticity,

- a centralized sovereign judge,
- a social scoring system,
- a psychological mind-reader,
- a post-hoc audit layer,
- or a replacement for ARL, L4 Witness, or hardware perimeter controls.

AGL is much narrower. It exists to prevent the system from acting as though a source is operationally grounded when it is not.

6. Relationship to the Existing Stack

6.1 Relation to ARL

ARL governs what happens once a dispute is admitted into procedural form: standing, admissibility, freeze, review, outcome, and re-entry discipline.

AGL sits **upstream** of that. It determines whether the initiating source is grounded enough to support runtime reliance before ordinary progression, escalation, or review begins.

6.2 Relation to L4 Witness

L4 Witness records and binds what happened. AGL helps determine whether the originating state was sufficiently grounded to authorize runtime dependence before that trace grows more costly downstream.

6.3 Relation to L4 Hardware Perimeter

Hardware perimeter controls define real operational preconditions: network mode, radios, power path, physical access, anomaly flags, and circuit-breaker behavior. AGL consumes this reality as grounding input instead of pretending execution begins from a neutral abstract plane.

7. Why the Canonical Home Is `ester-reality-bound`

AGL belongs canonically in `ester-reality-bound` because the problem it addresses is not merely argumentative or juridical. It is about **real present execution state**.

That includes:

- physical perimeter,
- operator condition,
- sensor trustworthiness,
- live anomaly state,
- and whether the system is acting under actual runtime grounding rather than procedural theater.

ARL consumes grounded dispute conditions. AGL helps decide whether those conditions are grounded enough to matter in the first place.

8. Why This Matters

A system can already be too late by the time review begins.

The structure may still look valid. The logs may still look coherent. The authority path may still look formally clean.

But if the initiating source was already degraded, detached, simulated, or contextually ungrounded, then the system is only reviewing the downstream shape of a bad beginning.

That is why grounding must become first-class. Not as a philosophical luxury, but as a condition for sane long-lived execution.

9. Explicit Bridge

L4 Hardware Perimeter <-> Actor Grounding Layer <-> ARL <-> L4 Witness

Grounding must exist before lawful reliance. Review must exist before lawful continuation. Witness must exist before trustworthy memory.

10. Hidden Bridges

Hidden bridge #1 - ERB <-> SER

AGL explains why **pre-admissibility hold** is not merely a procedural courtesy. It is the point where grounding failure can stop the path before ordinary review is forced to clean up a bad initiation.

Hidden bridge #2 - ERB <-> ECC

AGL gives the missing upstream meaning to implementation-facing concepts such as: intake hooks, hold-entry hooks, privilege narrowing, and runtime refusal. Without grounding discipline, those hooks remain mechanically useful but semantically incomplete.

11. Earth Paragraph

In a real operating room, no serious team says: “the scalpel is sterile, the checklist is complete, and the chart is coherent - so proceed,” without also asking whether the hand is steady, the patient is oxygenated, the monitors are live, and the room is still inside safe conditions.

If the source state is wrong, you stop **before** the cut. Not after the incision, not after the report, not after the review meeting.

AGL is that logic for long-lived digital systems.

12. Summary

The Actor Grounding Layer (AGL) formalizes a simple rule:

**before a system may rely on a source strongly enough to act, escalate,
or review,
it must determine whether that source is grounded enough in real
present execution state to count at all.**

This is not a replacement for review. It is the missing upstream condition that keeps review from becoming too late.